

Clustering Items through Bandit Feedback

Finding the Right Feature out of Many

Maximilian Graf* **Victor Thuot**[†] Nicolas Verzelen[†]

September 5th,  **StatMathAppli** 2025, Fréjus


* Institut für Mathematik, Universität Potsdam, Potsdam, Germany

[†] INRAE, Mistea, Institut Agro, Univ Montpellier, Montpellier, France

Setting

Motivating example

- forest patches, divided into two unknown groups
 - want to use biodiversity sensors to recover the groups
- ⇒ partition the forest by allocating efficiently the sensors
1. Which sensor is the most informative ?
 2. What budget (of observation) do we need ?



Clustering Items ...

- n : number of forest patches (items)
- d : number of sensors (features)
- $M_{i,j}$: mean value of the j -th sensor on the i -th patch

$$M = \begin{bmatrix} M_{1,1} & \cdots & M_{1,j} & \cdots & M_{1,d} \\ \vdots & & \vdots & & \vdots \\ M_{i,1} & \cdots & M_{i,j} & \cdots & M_{i,d} \\ \vdots & & \vdots & & \vdots \\ M_{n,1} & \cdots & M_{n,j} & \cdots & M_{n,d} \end{bmatrix} \leftarrow M_{i,\cdot} = \{m_{i,0}, m_{i,1}\}$$

Clustering Items ...

- n : number of forest patches (items)
- d : number of sensors (features)
- $M_{i,j}$: mean value of the j -th sensor on the i -th patch

$$M = \begin{bmatrix} M_{1,1} & \cdots & M_{1,j} & \cdots & M_{1,d} \\ \vdots & & \vdots & & \vdots \\ M_{i,1} & \cdots & M_{i,j} & \cdots & M_{i,d} \\ \vdots & & \vdots & & \vdots \\ M_{n,1} & \cdots & M_{n,j} & \cdots & M_{n,d} \end{bmatrix} \leftarrow M_{i,\cdot} = \{m_{i,0}, m_{i,1}\}$$

Clustering Items ...

- n : number of forest patches (items)
- d : number of sensors (features)
- $M_{i,j}$: mean value of the j -th sensor on the i -th patch

$$M = \begin{bmatrix} M_{1,1} & \cdots & M_{1,j} & \cdots & M_{1,d} \\ \vdots & & \vdots & & \vdots \\ M_{i,1} & \cdots & M_{i,j} & \cdots & M_{i,d} \\ \vdots & & \vdots & & \vdots \\ M_{n,1} & \cdots & M_{n,j} & \cdots & M_{n,d} \end{bmatrix} \leftarrow M_{i,\cdot} = \{m_{i,0}, m_{i,1}\}$$

Clustering Items ...

- n : number of forest patches (items)
- d : number of sensors (features)
- $M_{i,j}$: mean value of the j -th sensor on the i -th patch

$$M = \begin{bmatrix} M_{1,1} & \cdots & M_{1,j} & \cdots & M_{1,d} \\ \vdots & & \vdots & & \vdots \\ M_{i,1} & \cdots & M_{i,j} & \cdots & M_{i,d} \\ \vdots & & \vdots & & \vdots \\ M_{n,1} & \cdots & M_{n,j} & \cdots & M_{n,d} \end{bmatrix} \leftarrow M_{i,\cdot} \in \{\mu_0, \mu_1\}$$












- **Hidden partition**: there are **two groups** of patches

$$\boxed{\exists \mu_0 \neq \mu_1 \in \mathbb{R}^d; \forall i \in \{1, \dots, n\}, M_{i,\cdot} \in \{\mu_0, \mu_1\}}$$

- **Gap vector** Denote $\Delta = \mu_1 - \mu_0$ (w.l.o.g, $\mu_0 = 0$)

Clustering Items ...

- n : number of forest patches (items)
- d : number of sensors (features)

						
	0	0	0	0	0	0
	0	1	0	0.5	0.05	0
	0	1	0	0.5	0.05	0
	0	0	0	0	0	0
	0	1	0	0.5	0.05	0

- **Hidden partition**: there are **two groups** of patches

$$\exists \mu_0 \neq \mu_1 \in \mathbb{R}^d; \forall i \in \{1, \dots, n\}, M_{i,\cdot} \in \{\mu_0, \mu_1\}$$

- **Gap vector** Denote $\Delta = \mu_1 - \mu_0$ (w.l.o.g, $\mu_0 = 0$)

...through Bandit feedback

Vanilla clustering: observe the entire matrix

Bandit clustering: adapt the learning strategy on the fly

...through Bandit feedback

Vanilla clustering: observe the entire matrix

Bandit clustering: adapt the learning strategy on the fly

Learning protocol

At each time step t ,

- choose a patch $I_t \in 1, \dots, n$ (*based on the past*)
- choose a sensor $J_t \in 1, \dots, d$ (*based on the past*)
- receive X_t , s.t., $X_t = M_{I_t, J_t} + \text{random noise}$

Sampling protocol (example)

$$\begin{bmatrix} \boxed{0} & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

A red arrow points down to the top-left element (0) of the matrix, and another red arrow points left to the first row of the matrix.

Time t	1	2	3	4
(patch,sensor)	(1,1)			
$X_t = \boxed{} + \text{noise}$				

Sampling protocol (example)

$$\begin{bmatrix} \boxed{0} & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

A red arrow points down to the top-left element (0) of the matrix, and another red arrow points left to the first row of the matrix.

Time t	1	2	3	4
(patch,sensor)	(1,1)			
$X_t = \boxed{0.1} + \text{noise}$	0.1			

Sampling protocol (example)

$$\begin{array}{c} \downarrow \\ \left[\begin{array}{cccccc} 0 & 0 & 0 & 0 & 0 & 0 \\ \textcolor{yellow}{0} & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{array} \right] \leftarrow \end{array}$$

Time t	1	2	3	4
(patch,sensor)	(1,1)	(2,1)		
$X_t = \textcolor{yellow}{} + \text{noise}$	0.1			

Sampling protocol (example)

$$\begin{array}{c} \downarrow \\ \left[\begin{array}{cccccc} 0 & 0 & 0 & 0 & 0 & 0 \\ \textcolor{yellow}{0} & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{array} \right] \leftarrow \end{array}$$

Time t	1	2	3	4
(patch,sensor)	(1,1)	(2,1)		
$X_t = \textcolor{yellow}{} + \text{noise}$	0.1	-0.05		

Sampling protocol (example)

$$\begin{bmatrix} 0 & 0 & 0 & \boxed{0} & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

↓

←

Time t	1	2	3	4
(patch,sensor)	(1,1)	(2,1)	(1,4)	
$X_t = \boxed{} + \text{noise}$	0.1	-0.05		

Sampling protocol (example)

$$\begin{bmatrix} 0 & 0 & 0 & \boxed{0} & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

A red arrow points down to the highlighted '0' in the first row, fourth column. A red arrow points left to the closing bracket of the matrix.

Time t	1	2	3	4
(patch,sensor)	(1,1)	(2,1)	(1,4)	
$X_t = \boxed{} + \text{noise}$	0.1	-0.05	0.1	

Sampling protocol (example)

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

A red arrow points down to the top row, and a red arrow points left to the fourth column. The value 0.5 in the fourth row, fourth column is highlighted with a yellow box.

Time t	1	2	3	4
(patch,sensor)	(1,1)	(2,1)	(1,4)	(4,4)
$X_t = $ $+ \text{noise}$	0.1	-0.05	0.1	

Sampling protocol (example)


$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

A red arrow points down to the first '0' in the first row, fourth column. A red arrow points left to the '0.5' in the fourth row, fourth column, which is highlighted with a yellow background.

Time t	1	2	3	4
(patch,sensor)	(1,1)	(2,1)	(1,4)	(4,4)
$X_t = $ $+ \text{noise}$	0.1	-0.05	0.1	0.4

Sampling protocol (example)

$$\begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0.5 & 0.1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Time t	1	2	3	4	...
(patch,sensor)	(1,1)	(2,1)	(1,4)	(4,4)	
$X_t = $  $+ \text{noise}$	0.1	-0.05	0.1	0.4	...

→ At time T , output a partition of the patches

PAC-setting

Learning protocol

Input: prescribed probability δ

While $t \leq T$, (T a stopping time)

- choose a patch $I_t \in 1, \dots, n$
- choose a sensor $J_t \in 1, \dots, d$
- receive X_t , s.t., $X_t = M_{I_t, J_t} + \text{noise}^a$

Output: \hat{g} partition of the patches

a. $\epsilon_t := (X_t - M_{I_t, J_t})$ is 1-sub-Gaussian

PAC-setting

Learning protocol

Input: prescribed probability δ

While $t \leq T$, (T a stopping time)

- choose a patch $I_t \in 1, \dots, n$
- choose a sensor $J_t \in 1, \dots, d$
- receive X_t , s.t., $X_t = M_{I_t, J_t} + \text{noise}^a$

Output: \hat{g} partition of the patches

a. $\epsilon_t := (X_t - M_{I_t, J_t})$ is 1-sub-Gaussian

Definition

\mathcal{A} is δ -PAC if $\mathbb{P}_{\mathcal{A}, M}(\hat{g} \text{ is correct}) \geq 1 - \delta$

PAC-setting

Learning protocol

Input: prescribed probability δ

While $t \leq T$, (T a stopping time)

- choose a patch $I_t \in 1, \dots, n$
- choose a sensor $J_t \in 1, \dots, d$
- receive X_t , s.t., $X_t = M_{I_t, J_t} + \text{noise}^a$

Output: \hat{g} partition of the patches

a. $\epsilon_t := (X_t - M_{I_t, J_t})$ is 1-sub-Gaussian

Definition

\mathcal{A} is δ -PAC if $\mathbb{P}_{\mathcal{A}, M}(\hat{g} \text{ is correct}) \geq 1 - \delta$

Main objective

find \mathcal{A} , δ -PAC algorithm with a budget T as small as possible

Connection to pure exploration literature

- The problem¹ consists in balancing the exploration
 - $\left\{ \begin{array}{ll} \text{over the patches} & \rightarrow \text{to classify} \\ \text{over the sensors} & \rightarrow \text{to learn the structure of } \Delta \end{array} \right.$
- We combine ideas from (active) signal detection²³, and good-arm-identification⁴.

1. K. Ariu et al. “**Optimal clustering from noisy binary feedback**”. In: *Machine Learning* (2024).

2. R. M Castro. “**Adaptive sensing performance lower bounds for sparse signal detection and support estimation**”. In: *Bernoulli* (2014).

3. M. Saad, N. Verzelen, and A. Carpentier. “**Active ranking of experts based on their performances in many tasks**”. In: *ICML*. PMLR. 2023.

4. Y. Zhao et al. “**Revisiting simple regret: Fast rates for returning a good arm**”. In: *ICML*. PMLR. 2023.

Contribution

Our Contribution

1. Introduction of a δ -PAC Algorithm `BanditClustering`, which
 - a) identifies a sensor discriminative enough
 - b) classify patches based on this sensor

Our Contribution

1. Introduction of a δ -PAC Algorithm `BanditClustering`, which
 - a) identifies a sensor discriminative enough
 - b) classify patches based on this sensor
2. Answers to the questions:
 - What is the sampling budget required ?
(upper bound on the budget of `BanditClustering`)
 - Is our approach optimal?
(matching lower in most relevant cases)

Our Contribution

1. Introduction of a δ -PAC Algorithm `BanditClustering`, which
 - a) identifies a sensor discriminative enough
 - b) classify patches based on this sensor
2. Answers to the questions:
 - What is the sampling budget required ?
(upper bound on the budget of `BanditClustering`)
 - Is our approach optimal?
(matching lower in most relevant cases)
3. Illustrative numerical experiments

Algorithm: BanditClustering

1. Identification of **representatives patches** r_0, r_1
(\sim signal detection)

$$\Delta = M_{r_1, \cdot} - M_{r_0, \cdot}$$

2. Selection of j a good **discriminative sensor**
(\sim good-arm identification)

$$\Delta_j = M_{r_1, j} - M_{r_0, j}$$

3. **Classification** of patches based on the sensor
(\sim binary classification)

First step : representative identification

- **Objective:** For r_0 fixed, find r_1 s.t. $M_{r_1,\cdot} \neq M_{r_0,\cdot}$
(w.l.o.g, $M_{r_0,\cdot} = 0$, and $M_{r_1,\cdot} = \Delta$)

First step : representative identification

- **Objective:** For r_0 fixed, find r_1 s.t. $M_{r_1,\cdot} \neq M_{r_0,\cdot}$
(w.l.o.g, $M_{r_0,\cdot} = 0$, and $M_{r_1,\cdot} = \Delta$)
- **Method:**
 1. **subsampling:** select randomly L entries of M
 2. **Sequential halving:** identify entry (i,j) with $M_{i,j}$ large (with SH and budget N)
 3. **stopping condition:** if $M_{i,j} \neq 0$ w.h.p
output the patch of this entry
Else increase (L, N) and repeat 1, 2, 3

First step : representative identification

- **Objective:** For r_0 fixed, find r_1 s.t. $M_{r_1,\cdot} \neq M_{r_0,\cdot}$
(w.l.o.g, $M_{r_0,\cdot} = 0$, and $M_{r_1,\cdot} = \Delta$)
- **Method:**
 1. **subsampling:** select randomly L entries of M
 2. **Sequential halving:** identify entry (i,j) with $M_{i,j}$ large (with SH and budget N)
 3. **stopping condition:** if $M_{i,j} \neq 0$ w.h.p
output the patch of this entry
Else increase (L, N) and repeat 1, 2, 3
- **Budget:**

$$\frac{d}{\theta \|\Delta\|_2^2} \log(1/\delta) \quad (\text{up to log terms})$$

with θ proportion of patches in the smallest group

Second step : sensor selection

- **Objective:** We know r_1 s.t. $\Delta = M_{r_1, \cdot} \neq M_{r_0, \cdot} = 0$, we want j (sensor) s.t. Δ_j is large enough

Second step : sensor selection

- **Objective:** We know r_1 s.t. $\Delta = M_{r_1, \cdot} \neq M_{r_0, \cdot} = 0$, we want j (sensor) s.t. Δ_j is large enough
- **Method:**
 1. **subsampling:** select randomly L sensors among $\{1, \dots, d\}$
 2. **Sequential halving:** identify entry j with Δ_j large (with SH and budget N)
 3. **Estimation:** estimate Δ_j
 4. **stopping condition:** if $\frac{n}{\hat{\Delta}_j^2} \log(n/\delta) \lesssim N$ w.h.p

classification

Else increase (L, N) and repeat 1, 2, 3

Second step (and third step): sensor selection

- **Objective:** We know r_1 s.t. $\Delta = M_{r_1, \cdot} \neq M_{r_0, \cdot} = 0$, we want j (sensor) s.t. Δ_j is large enough
- **Method:**
 1. **subsampling:** select randomly L sensors among $\{1, \dots, d\}$
 2. **Sequential halving:** identify entry j with Δ_j large (with SH and budget N)
 3. **Estimation:** estimate Δ_j
 4. **stopping condition:** if $\frac{n}{\Delta_j^2} \log(n/\delta) \lesssim N$ w.h.p

Sample each patch i with sensor j , N/n times

Use a linear classifier

Else increase (L, N) and repeat 1, 2, 3

Theorem (general case)

Theorem

Define

$$H := \frac{d}{\theta} \left(\frac{1}{\|\Delta\|^2} \right) + \min_{s \in [d]} \left(\frac{d}{s} + n \right) \left(\frac{1}{\Delta_{(s)}^2} \right) ,$$

with $|\Delta_{(1)}| \geq |\Delta_{(2)}| \geq \dots$

With probability $1 - \delta$, BanditClustering returns the good partition after

$$T \lesssim \log \left(\frac{1}{\delta} \right) \cdot H$$

Theorem (Simpler case)

recall: matrix $(M_{i,j})_{\substack{i=1,\dots,n \\ j=1,\dots,d}}$ with rows $M_{i,\cdot} \in \{0, \Delta\}$

Theorem (Simpler case)

recall: matrix $(M_{i,j})_{\substack{i=1,\dots,n \\ j=1,\dots,d}}$ with rows $M_{i,\cdot} \in \{0, \Delta\}$

sparse setting: Δ equals $h > 0$ in s entries and 0 elsewhere

Theorem (Simpler case)

recall: matrix $(M_{i,j})_{\substack{i=1,\dots,n \\ j=1,\dots,d}}$ with rows $M_{i,\cdot} \in \{0, \Delta\}$

sparse setting: Δ equals $h > 0$ in s entries and 0 elsewhere

Theorem

With probability $1 - \delta$ BanditClustering returns the good partition after $T \lesssim \log(1/\delta) \cdot H$, where

$$H := \frac{1}{\theta} \frac{d}{\|\Delta\|^2} + \frac{n}{h^2} .$$

Theorem

If \mathcal{A} is δ -PAC, then there exists a permutation of M , M_{per} , s.t.,

$$\mathbb{P}_{M_{per}, \mathcal{A}} \left(T \geq \frac{2d}{\theta \|\Delta\|^2} \log(1/6\delta) \vee \frac{2(n-2)}{\|\Delta\|_\infty^2} \log(1/4.8\delta) \right) \geq \delta .$$

Theorem

If \mathcal{A} is δ -PAC, then there exists a permutation of M , M_{per} , s.t.,

$$\mathbb{P}_{M_{per}, \mathcal{A}} \left(T \geq \frac{2d}{\theta \|\Delta\|^2} \log(1/6\delta) \vee \frac{2(n-2)}{\|\Delta\|_\infty^2} \log(1/4.8\delta) \right) \geq \delta .$$

Assume that $\Delta = (\underbrace{h, \dots, h}_s, 0, \dots, 0)$

Lower bound

Theorem

If \mathcal{A} is δ -PAC, then there exists a permutation of M , M_{per} , s.t.,

$$\mathbb{P}_{M_{per}, \mathcal{A}} \left(T \geq \frac{2d}{\theta \|\Delta\|^2} \log(1/6\delta) \vee \frac{2(n-2)}{\|\Delta\|_\infty^2} \log(1/4.8\delta) \right) \geq \delta .$$

Assume that $\Delta = (\underbrace{h, \dots, h}_s, 0, \dots, 0)$

1. Finding a patch in the smallest group \rightarrow explore $\frac{1}{\theta} \log(1/\delta)$

Theorem

If \mathcal{A} is δ -PAC, then there exists a permutation of M , M_{per} , s.t.,

$$\mathbb{P}_{M_{per}, \mathcal{A}} \left(T \geq \frac{2d}{\theta \|\Delta\|^2} \log(1/6\delta) \vee \frac{2(n-2)}{\|\Delta\|_\infty^2} \log(1/4.8\delta) \right) \geq \delta .$$

Assume that $\Delta = (\underbrace{h, \dots, h}_s, 0, \dots, 0)$

1. Finding a patch in the smallest group \rightarrow explore $\frac{1}{\theta} \log(1/\delta)$
2. Detecting a sensor j with $\Delta_j = h$ \rightarrow explore $\frac{d}{s} \log(1/\delta)$

Lower bound

Theorem

If \mathcal{A} is δ -PAC, then there exists a permutation of M , M_{per} , s.t.,

$$\mathbb{P}_{M_{per}, \mathcal{A}} \left(T \geq \frac{2d}{\theta \|\Delta\|^2} \log(1/6\delta) \vee \frac{2(n-2)}{\|\Delta\|_\infty^2} \log(1/4.8\delta) \right) \geq \delta .$$

Assume that $\Delta = (\underbrace{h, \dots, h}_s, 0, \dots, 0)$

1. Finding a patch in the smallest group \rightarrow explore $\frac{1}{\theta} \log(1/\delta)$
2. Detecting a sensor j with $\Delta_j = h$ \rightarrow explore $\frac{d}{s} \log(1/\delta)$
3. Testing $\Delta_j = 0$ VS $\Delta_j = h$ \rightarrow sample $\frac{1}{h^2} \log(1/\delta)$

Theorem

If \mathcal{A} is δ -PAC, then there exists a permutation of M , M_{per} , s.t.,

$$\mathbb{P}_{M_{per}, \mathcal{A}} \left(T \geq \frac{2d}{\theta \|\Delta\|^2} \log(1/6\delta) \vee \frac{2(n-2)}{\|\Delta\|_\infty^2} \log(1/4.8\delta) \right) \geq \delta .$$

Theorem

If \mathcal{A} is δ -PAC, then there exists a permutation of M , M_{per} , s.t.,

$$\mathbb{P}_{M_{per}, \mathcal{A}} \left(T \geq \frac{2d}{\theta \|\Delta\|^2} \log(1/6\delta) \vee \frac{2(n-2)}{\|\Delta\|_\infty^2} \log(1/4.8\delta) \right) \geq \delta .$$

- Let j such that Δ_j is maximal, for each item, we need at least $\frac{1}{\Delta_j^2} \log(1/\delta)$ samples from to classify it

Theorem

If \mathcal{A} is δ -PAC, then there exists a permutation of M , M_{per} , s.t.,

$$\mathbb{P}_{M_{per}, \mathcal{A}} \left(T \geq \frac{2d}{\theta \|\Delta\|^2} \log(1/6\delta) \vee \frac{2(n-2)}{\|\Delta\|_\infty^2} \log(1/4.8\delta) \right) \geq \delta .$$

- Let j such that Δ_j is maximal, for each item, we need at least $\frac{1}{\Delta_j^2} \log(1/\delta)$ samples from to classify it
- **Remark.** Still a gap with $\min_{s \in [d]} \left(\frac{d}{s} + n \right) \frac{1}{\Delta_{(s)}^2}$, however, optimal if Δ takes two values.

Numerical experiments




Figure 1: Budgets for increasing sparsity s with fixed $\|\Delta\|$
 $n = 20, d = 1000, \Delta^s = (\underbrace{h_s, \dots, h_s}_{s \text{ times}}, 0, \dots, 0)$, and $h_s = \frac{15}{\sqrt{s}}$

1. **Algorithmic solution:** new online clustering alg. of budget

$$\frac{d}{\theta} \left(\frac{1}{\|\Delta\|^2} \right) \log \left(\frac{1}{\delta} \right) + \min_{s \in [d]} \left(\frac{d}{s} + n \right) \left(\frac{1}{\Delta_{(s)}^2} \right) \log \left(\frac{1}{\delta} \right)$$

Take home message

1. **Algorithmic solution:** new online clustering alg. of budget

$$\frac{d}{\theta} \left(\frac{1}{\|\Delta\|^2} \right) \log \left(\frac{1}{\delta} \right) + \min_{s \in [d]} \left(\frac{d}{s} + n \right) \left(\frac{1}{\Delta_{(s)}^2} \right) \log \left(\frac{1}{\delta} \right)$$

2. **Optimality:** matching lower bound in the sparse setting

Take home message

1. **Algorithmic solution:** new online clustering alg. of budget

$$\frac{d}{\theta} \left(\frac{1}{\|\Delta\|^2} \right) \log \left(\frac{1}{\delta} \right) + \min_{s \in [d]} \left(\frac{d}{s} + n \right) \left(\frac{1}{\Delta_{(s)}^2} \right) \log \left(\frac{1}{\delta} \right)$$

2. **Optimality:** matching lower bound in the sparse setting
3. **Perspectives:** refine the lower bounds, generalize to $K > 2$

Take home message

1. **Algorithmic solution:** new online clustering alg. of budget

$$\frac{d}{\theta} \left(\frac{1}{\|\Delta\|^2} \right) \log \left(\frac{1}{\delta} \right) + \min_{s \in [d]} \left(\frac{d}{s} + n \right) \left(\frac{1}{\Delta_{(s)}^2} \right) \log \left(\frac{1}{\delta} \right)$$

2. **Optimality:** matching lower bound in the sparse setting
3. **Perspectives:** refine the lower bounds, generalize to $K > 2$
4. **Reference:**

M. Graf, V. Thuot, and N. Verzelen. *Clustering Items through Bandit Feedback: Finding the Right Feature out of Many*. Proceedings of the 42nd International Conference on Machine Learning. 2025. arXiv: 2503.11209

Thank you

Sequential Halving

- 1: **Input:** Number of arms M , budget T
- 2: Set $S \leftarrow \{1, 2, \dots, M\}$ ▷ Set of active arms
- 3: Set $R \leftarrow \lfloor \log_2 M \rfloor$ ▷ Number of rounds
- 4: **for** $r = 1$ to R **do**
- 5: $n_r \leftarrow \left\lfloor \frac{T}{|S| \cdot R} \right\rfloor$ ▷ $n_{r+1} = 2 \cdot n_r$
- 6: **for** each arm a in S **do**
- 7: Sample arm a for n_r times, compute empirical mean $\hat{\mu}_a$
- 8: **end for**
- 9: Keep top $\left\lfloor \frac{|S|}{2} \right\rfloor$ arms in S
- 10: **end for**[final]beamer
- 11: **Output:** Return the remaining arm in S